

Pattern Matching

Document #: D1260R1
Date: 2018-11-09
Project: Programming Language C++
Evolution
Reply-to: Michael Park
<mcypark@gmail.com>
David Sankel
<dsankel@bloomberg.net>
Dan Sarginson
<dsarginson@bloomberg.net>
Sergei Murzin
<smurzin@bloomberg.net>

Contents

1	Revision History	2
2	Introduction	2
3	Motivation and Scope	3
4	Before/After Comparisons	3
4.1	Matching Integrals	3
4.2	Matching Strings	4
4.3	Matching Tuples	4
4.4	Matching Variants	4
4.5	Matching Polymorphic Types	5
4.6	Evaluating Expression Trees	5
5	Design Overview	6
5.1	Basic Syntax	6
5.2	Basic Model	7
5.3	Types of Patterns	7
5.3.1	Primary Patterns	7
5.3.2	Compound Patterns	8
5.4	Pattern Guard	13
5.5	<code>inspect constexpr</code>	14
5.6	Exhaustiveness and Usefulness	14
6	Proposed Wording	15
7	Design Decisions	15
7.1	Extending Structured Bindings Declaration	15
7.2	<code>inspect</code> rather than <code>switch</code>	16
7.3	First Match rather than Best Match	16
7.4	Language rather than Library	16

7.5	Matchers and Extractors	16
7.6	Wildcard Syntax	17
8	Runtime Performance	17
8.0.1	Structured Binding Pattern	17
8.0.2	Alternative Pattern	17
8.0.3	Open Class Hierarchy	17
9	Examples	17
9.1	Predicate-based Discriminator	17
9.2	“Closed” Class Hierarchy	19
9.3	Matcher: <code>any_of</code>	20
9.4	Matcher: <code>within</code>	20
9.5	Extractor: <code>both</code>	21
9.6	Extractor: <code>at</code>	21
10	Future Work	21
10.1	Language Support for Variant	21
10.2	Patterns in range-based <code>for</code> loop	22
10.3	Note on Ranges	23
11	Acknowledgements	23
	References	23

1 Revision History

- R1
 - Added an [Examples](#) section
 - Strong support for pattern matching in EWGI, San Diego.
 “Should we commit additional committee time to pattern matching?”
 SF: 14, WF: 0, N: 1, WA: 0, SA: 0
 - Added the expression form of `inspect`
 - Added the dereference (`*`) pattern
 - Added the designated (`.field`) extractors
 - Added the pin (`^`) operator for the expression pattern
 - Used `<>` for `std::any`-like and polymorphic types
- R0 — Initial Draft

2 Introduction

As algebraic data types gain better support in C++ with facilities such as `tuple` and `variant`, the importance of mechanisms to interact with them have increased. While mechanisms such as `apply` and `visit` have been added, their usage is quite complex and limited even for simple cases. Pattern matching is a widely adopted mechanism across many programming languages to interact with algebraic data types that can help greatly simplify C++. Examples of programming languages include text-based languages such as SNOBOL back in the 1960s, functional languages such as Haskell and OCaml, and “mainstream” languages such as Scala, Swift, and Rust.

Inspired by P0095 [5] — which proposed pattern matching and language-level variant simultaneously — this paper explores a possible direction for pattern matching only, and does not address language-level variant design. This is in correspondence with a straw poll from Kona 2015, which encouraged exploration of a full solution for pattern matching. SF: 16, WF: 6, N: 5, WA: 1, SA: 0.

3 Motivation and Scope

Virtually every program involves branching on some predicates applied to a value and conditionally binding names to some of its components for use in subsequent logic. Today, C++ provides two types of selection statements: the `if` statement and the `switch` statement.

Since `switch` statements can only operate on a *single* integral value and `if` statements operate on an *arbitrarily* complex boolean expression, there is a significant gap between the two constructs even in inspection of the “vocabulary types” provided by the standard library.

In C++17, structured binding declarations [9] introduced the ability to concisely bind names to components of `tuple`-like values. The proposed direction of this paper aims to naturally extend this notion by performing **structured inspection** prior to forming the **structured bindings** with a third selection statement: the `inspect` statement. The goal of the `inspect` statement is to bridge the gap between `switch` and `if` statements with a **declarative**, **structured**, **cohesive**, and **composable** mechanism.

4 Before/After Comparisons

4.1 Matching Integrals

Before

```
switch (x) {
  case 0: std::cout << "got zero"; break;
  case 1: std::cout << "got one"; break;
  default: std::cout << "don't care";
}
```

After

```
inspect (x) {
  0: std::cout << "got zero";
  1: std::cout << "got one";
  _: std::cout << "don't care";
}
```

4.2 Matching Strings

Before

```
if (s == "foo") {
    std::cout << "got foo";
} else if (s == "bar") {
    std::cout << "got bar";
} else {
    std::cout << "don't care";
}
```

After

```
inspect (s) {
    "foo": std::cout << "got foo";
    "bar": std::cout << "got bar";
    _: std::cout << "don't care";
}
```

4.3 Matching Tuples

Before

```
auto&& [x, y] = p;
if (x == 0 && y == 0) {
    std::cout << "on origin";
} else if (x == 0) {
    std::cout << "on y-axis";
} else if (y == 0) {
    std::cout << "on x-axis";
} else {
    std::cout << x << ', ' << y;
}
```

After

```
inspect (p) {
    [0, 0]: std::cout << "on origin";
    [0, y]: std::cout << "on y-axis";
    [x, 0]: std::cout << "on x-axis";
    [x, y]: std::cout << x << ', ' << y;
}
```

4.4 Matching Variants

Before

```
struct visitor {
    void operator()(int i) const {
        os << "got int: " << i;
    }
    void operator()(float f) const {
        os << "got float: " << f;
    }
    std::ostream& os;
};
std::visit(visitor{strm}, v);
```

After

```
inspect (v) {
    <int> i: strm << "got int: " << i;
    <float> f: strm << "got float: " << f;
}
```

4.5 Matching Polymorphic Types

```
struct Shape { virtual ~Shape() = default; };
struct Circle : Shape { int radius; };
struct Rectangle : Shape { int width, height; };
```

Before

```
virtual int Shape::get_area() const = 0;

int Circle::get_area() const override {
    return 3.14 * radius * radius;
}
int Rectangle::get_area() const override {
    return width * height;
}
```

After

```
int get_area(const Shape& shape) {
    inspect (shape) {
        <Circle> [r]: return 3.14 * r * r;
        <Rectangle> [w, h]: return w * h;
    }
}
```

4.6 Evaluating Expression Trees

```
struct Expr;
struct Neg { std::shared_ptr<Expr> expr; };
struct Add { std::shared_ptr<Expr> lhs, rhs; };
struct Mul { std::shared_ptr<Expr> lhs, rhs; };
struct Expr : std::variant<int, Neg, Add, Mul> { using variant::variant; };

namespace std {
    template <>
    struct variant_size<Expr> : variant_size<Expr::variant> {};

    template <std::size_t I>
    struct variant_alternative<I, Expr> : variant_alternative<I, Expr::variant> {};
}
```

Before

After

```
int eval(const Expr& expr) {
    struct visitor {
        int operator()(int i) const {
            return i;
        }
        int operator()(const Neg& n) const {
            return -eval(*n.expr);
        }
        int operator()(const Add& a) const {
            return eval(*a.lhs) + eval(*a.rhs);
        }
        int operator()(const Mul& m) const {
            return eval(*m.lhs) * eval(*m.rhs);
        }
    };
    return std::visit(visitor{}, expr);
}
```

```
int eval(const Expr& expr) {
    inspect (expr) {
        <int> i: return i;
        <Neg> [e]: return -eval(*e);
        <Add> [l, r]: return eval(*l) + eval(*r);
        <Mul> [l, r]: return eval(*l) * eval(*r);
    }
}
```

5 Design Overview

5.1 Basic Syntax

There are two forms of `inspect`: the statement form and the expression form.

```
inspect constexpropt ( init-statementopt condition ) {
    pattern guardopt : statement
    pattern guardopt : statement
    ...
}

inspect constexpropt ( init-statementopt condition ) trailing-return-typeopt {
    pattern guardopt => expression ,
    pattern guardopt => expression ,
    ...
}

guard:
    if ( expression )
```

[*Note:* The expression form is equivalent to:

```
std::invoke([&]() trailing-return-typeopt {
    inspect constexpropt ( init-statementopt condition ) {
        pattern guardopt : return expression ;
        pattern guardopt : return expression ;
        ...
    }
})
```

— *end note*]

5.2 Basic Model

Within the parentheses, the `inspect` statement is equivalent to `switch` and `if` statements except that no conversion nor promotion takes place in evaluating the value of its condition.

When the `inspect` statement is executed, its condition is evaluated and matched in order (first match semantics) against each pattern. If a pattern successfully matches the value of the condition and the boolean expression in the guard evaluates to `true` (or if there is no guard at all), control is passed to the statement following the matched pattern label. If the guard expression evaluates to `false`, control flows to the subsequent pattern.

If no pattern matches, none of the statements are executed for the statement form and `std::no_match` exception is thrown for the expression form.

5.3 Types of Patterns

5.3.1 Primary Patterns

5.3.1.1 Wildcard Pattern

The wildcard pattern has the form:

```
-  
and matches any value v.  
int v = /* ... */;  
  
inspect (v) {  
    _: std::cout << "ignored";  
    // ^ wildcard pattern  
}
```

[*Note*: Even though `_` is a valid identifier, it does not introduce a name. — *end note*]

5.3.1.2 Identifier Pattern

The identifier pattern has the form:

identifier

and matches any value `v`. The introduced name behaves as an lvalue referring to `v`, and is in scope from its point of declaration until the end of the statement following the pattern label.

```
int v = /* ... */;  
  
inspect (v) {  
    x: std::cout << x;  
    // ^ identifier pattern  
}
```

[*Note*: If the identifier pattern is used at the top-level, it has the same syntax as a `goto` label. — *end note*]

5.3.1.3 Expression Pattern

The expression pattern has the form:

literal
^ primary-expression

and matches value *v* if a call to member *e*.match(*v*) or else a non-member ADL-only match(*e*, *v*) is contextually convertible to bool and evaluates to true where *e* is the *literal* or *primary-expression*.

The default behavior of match(*x*, *y*) is *x* == *y*.

```
int v = /* ... */;

inspect (v) {
    0: std::cout << "got zero";
    1: std::cout << "got one";
    // ^ expression pattern
}

static constexpr int zero = 0, one = 1;
int v = /* ... */;

inspect (v) {
    ^zero: std::cout << "got zero";
    // ^^^^^ expression pattern
}
```

5.3.2 Compound Patterns

5.3.2.1 Structured Binding Pattern

The structured binding pattern has the following two forms:

[*pattern*₀ , *pattern*₁ , ... , *pattern*_N]
[*designator*₀ : *pattern*₀ , *designator*₁ : *pattern*₁ , ... , *designator*_N : *pattern*_N]

The first form matches value *v* if each *pattern*_{*i*} matches the *i*th component of *v*. The components of *v* are given by the structured binding declaration: auto&& [__e₀, __e₁, ..., __e_N] = *v*; where each __e_{*i*} are unique exposition-only identifiers.

```
std::pair<int, int> p = /* ... */;

inspect (p) {
    [0, 0]: std::cout << "on origin";
    [0, y]: std::cout << "on y-axis";
    // ^ identifier pattern
    [x, 0]: std::cout << "on x-axis";
    // ^ expression pattern
    [x, y]: std::cout << x << ', ' << y;
    // ^^^^^ structured binding pattern
}
```

The second form matches value *v* if each *pattern*_{*i*} matches the direct non-static data member of *v* named *identifier* from each *designator*_{*i*}. If an *identifier* from any *designator*_{*i*} does not refer to a direct non-static data member of *v*, the program is ill-formed.


```

struct Player { std::string name; int hitpoints; int coins; };

void get_hint(const Player& p) {
    inspect (p) {
        [.hitpoints: 1]: std::cout << "You're almost destroyed. Give up!\n";
        [.hitpoints: 10, .coins: 10]: std::cout << "I need the hints from you!\n";
        [.coins: 10]: std::cout << "Get more hitpoints!\n";
        [.hitpoints: 10]: std::cout << "Get more ammo!\n";
        [.name: n]: {
            if (n != "The Bruce Dickenson") {
                std::cout << "Get more hitpoints and ammo!\n";
            } else {
                std::cout << "More cowbell!\n";
            }
        }
    }
}

```

[*Note*: Unlike designated initializers, the order of the designators need not be the same as the declaration order of the members of the class. — *end note*]

5.3.2.2 Alternative Pattern

The alternative pattern has the following forms:

```

< auto > pattern
< concept > pattern
< type > pattern
< constant-expression > pattern

```

Let *v* be the value being matched and *V* be `std::remove_cvref_t<decltype(v)>`.

Let *Alt* be the entity inside the angle brackets.

Case 1: `std::variant-like`

If `std::variant_size_v<V>` is well-formed and evaluates to an integral, the alternative pattern matches *v* if *Alt* is compatible with the current index of *v* and *pattern* matches the active alternative of *v*.

Let *I* be the current index of *v* given by a member *v.index()* or else a non-member ADL-only *index(v)*. The active alternative of *v* is given by `std::variant_alternative_t<I, V>&` initialized by a member *v.get<I>()* or else a non-member ADL-only *get<I>(v)*.

Alt is compatible with *I* if one of the following four cases is true:

- *Alt* is `auto`
- *Alt* is a *concept* and `std::variant_alternative_t<I, V>` satisfies the *concept*.
- *Alt* is a *type* and `std::is_same_v<Alt, std::variant_alternative_t<I, V>>` is `true`
- *Alt* is a *constant-expression* that can be used in a `switch` and is the same value as *I*.

Before	After
<pre>std::visit([&](auto&& x) { strm << "got auto: " << x; }, v);</pre>	<pre>inspect (v) { <auto> x: strm << "got auto: " << x; }</pre>
<pre>std::visit([&](auto&& x) { using X = std::remove_cvref_t<decltype(x)>; if constexpr (C1<X>()) { strm << "got C1: " << x; } else if constexpr (C2<X>()) { strm << "got C2: " << x; } }, v);</pre>	<pre>inspect (v) { <C1> c1: strm << "got C1: " << c1; <C2> c2: strm << "got C2: " << c2; }</pre>
<pre>std::visit([&](auto&& x) { using X = std::remove_cvref_t<decltype(x)>; if constexpr (std::is_same_v<int, X>) { strm << "got int: " << x; } else if constexpr (std::is_same_v<float, X>) { strm << "got float: " << x; } }, v);</pre>	<pre>inspect (v) { <int> i: strm << "got int: " << i; <float> f: strm << "got float: " << f; }</pre>
<pre>std::variant<int, int> v = /* ... */; std::visit([&](int x) { strm << "got int: " << x; }, v);</pre>	<pre>std::variant<int, int> v = /* ... */; inspect (v) { <int> x: strm << "got int: " << x; }</pre>

Before	After
<pre> std::variant<int, int> v = /* ... */; std::visit([&](auto&& x) { switch (v.index()) { case 0: { strm << "got first: " << x; break; } case 1: { strm << "got second: " << x; break; } } }, v); </pre>	<pre> std::variant<int, int> v = /* ... */; inspect (v) { <0> x: strm << "got first: " << x; <1> x: strm << "got second: " << x; } </pre>

Case 2: std::any-like

< type > pattern

If *Alt* is a *type* and there exists a valid non-member ADL-only `any_cast<Alt>(&v)`, let *p* be its result. The alternative pattern matches if *p* contextually converted to `bool` evaluates to `true`, and *pattern* matches **p*.

Before	After
<pre> std::any a = 42; if (int* i = any_cast<int>(&a)) { std::cout << "got int: " << *i; } else if (float* f = any_cast<float>(&a)) { std::cout << "got float: " << *f; } </pre>	<pre> std::any a = 42; inspect (a) { <int> i: std::cout << "got int: " << i; <float> f: std::cout << "got float: " << f; } </pre>

Case 3: Polymorphic Types

< type > pattern

If *Alt* is a *type* and `std::is_polymorphic_v<V>` is `true`, let *p* be `dynamic_cast<Alt'*>(&v)` where *Alt'* has the same *cv*-qualifications as `decltype(&v)`. The alternative pattern matches if *p* contextually converted to `bool` evaluates to `true`, and *pattern* matches **p*.

While the **semantics** of the pattern is specified in terms of `dynamic_cast`, N3449 [8] describes techniques involving vtable pointer caching and hash conflict minimization that are implemented in the Mach7 [6] library, as well as mentions of further opportunities available for a compiler intrinsic.

Given the following definition of a **Shape** class hierarchy:

```
struct Shape { virtual ~Shape() = default; };

struct Circle : Shape { int radius; };
struct Rectangle : Shape { int width, height; };
```

Before

```
virtual int Shape::get_area() const = 0;

int Circle::get_area() const override {
    return 3.14 * radius * radius;

int Rectangle::get_area() const override {
    return width * height;
}
```

After

```
int get_area(const Shape& shape) {
    inspect (shape) {
        <Circle> [r]: return 3.14 * r * r;
        <Rectangle> [w, h]: return w * h;
    }
}
```

5.3.2.3 Binding Pattern

The binding pattern has the form:

identifier @ pattern

and matches value *v* if *pattern* matches it. The introduced name behaves as an lvalue referring to *v*, and is in scope from its point of declaration until the end of the statement following the pattern label.

```
std::variant<Point, /* ... */> v = /* ... */;

inspect (v) {
    <Point> p @ [x, y]: // ...
    //          ^^^^^^^^^ binding pattern
}
```

5.3.2.4 Dereference Pattern

The dereference pattern has the form:

**pattern*

and matches value *v* if *pattern* matches **v*.

```
struct Node {
    int value;
    std::unique_ptr<Node> lhs, rhs;
};

template <typename Visitor>
void print_leftmost(const Node& node) {
    inspect (node) {
        [.value: v, .lhs: nullptr]: std::cout << v << '\n';
        [.lhs: *l]: print_leftmost(l);
    }
}
```

```

    }
}

```

[*Note*: The dereference operation is performed unconditionally. — *end note*]

5.3.2.5 Extractor Pattern

The extractor pattern has the following two forms:

```

( constant-expression ! pattern )
( constant-expression ? pattern )

```

Let *c* be the *constant-expression*. The first form matches value *v* if *pattern* matches *e* where *e* is the result of a call to member *c.extract(v)* or else a non-member ADL-only *extract(c, v)*.

For second form, let *e* be the result of a call to member *c.try_extract(v)* or else a non-member ADL-only *try_extract(c, v)*. It matches value *v* if *e* is contextually convertible to *bool*, evaluates to *true*, and *pattern* matches **e*.

```

struct {
    std::optional<std::array<std::string_view, 2>>
    try_extract(std::string_view sv) const;
} email;

struct {
    std::optional<std::array<std::string_view, 3>>
    try_extract(std::string_view sv) const;
} phone_number;

inspect (s) {
    (email? [address, domain]): std::cout << "got an email";
    (phone_number? ["415", _, _]): std::cout << "got a San Francisco phone number";
    // ~~~~~ extractor pattern
}

```

5.4 Pattern Guard

The pattern guard has the form:

```

if ( expression )

```

Let *e* be the result of *expression* contextually converted to *bool*. If *e* is *true*, control is passed to the corresponding statement. Otherwise, control flows to the subsequent pattern.

The pattern guard allows to perform complex tests that cannot be performed within the *pattern*. For example, performing tests across multiple bindings:

```

inspect (p) {
    [x, y] if (test(x, y)): std::cout << x << ',' << y << " passed";
    // ~~~~~ pattern guard
}

```

This also diminishes the desire for fall-through semantics within the statements, an unpopular feature even in *switch* statements. For the reified semantics of the pattern guard, consider the following snippet:

```
switch (x) {
  case c1: if (cond1) { stmt1; break; } [[fallthrough]]
  case c2: if (cond2) { stmt2; break; } [[fallthrough]]
}
```

5.5 inspect constexpr

Every *pattern* is able to determine whether it matches value *v* as a boolean expression in isolation. Let *MATCHES* be the condition for which a *pattern* matches a value *v*. Ignoring any potential optimization opportunities, we're able to perform the following transformation:

inspect	if
<pre>inspect (v) { pattern1 if (cond1): stmt1 pattern2: stmt2 // ... }</pre>	<pre>if (MATCHES(pattern1, v) && cond1) stmt1 else if (MATCHES(pattern2, v)) stmt2 // ...</pre>

`inspect constexpr` is then formulated by applying `constexpr` to every `if` branch.

inspect constexpr	if constexpr
<pre>inspect constexpr (v) { pattern1 if (cond1): stmt1 pattern2: stmt2 // ... }</pre>	<pre>if constexpr (MATCHES(pattern1, v) && cond1) stmt1 else if constexpr (MATCHES(pattern2, v)) stmt2 // ...</pre>

5.6 Exhaustiveness and Usefulness

`inspect` can be declared `[[strict]]` for implementation-defined exhaustiveness and usefulness checking.

Exhaustiveness means that all values of the type of the value being matched is handled by at least one of the cases. For example, having a `_:` case makes any `inspect` statement exhaustive.

Usefulness means that every case handles at least one value of the type of the value being matched. For example, any case that comes after a `_:` case would be useless.

Warnings for pattern matching [2] discusses and outlines an algorithm for exhaustiveness and usefulness for OCaml, and is the algorithm used by Rust.

6 Proposed Wording

The following is the beginning of an attempt at a syntactic structure.

Add to §8.4 [stmt.select] of ...

- ¹ Selection statements choose one of several flows of control.

selection-statement:

```
if constexpropt ( init-statementopt condition ) statement
if constexpropt ( init-statementopt condition ) statement else statement
switch ( init-statementopt condition ) statement
inspect constexpropt ( init-statementopt condition ) trailing-return-typeopt { inspect-case-seq }
```

inspect-case-seq:

```
inspect-statement-case-seq
inspect-expression-case-seq
```

inspect-statement-case-seq:

```
inspect-statement-case
inspect-statement-case-seq inspect-statement-case
```

inspect-expression-case-seq:

```
inspect-expression-case
inspect-expression-case-seq , inspect-expression-case
```

inspect-statement-case:

```
inspect-pattern inspect-guardopt : statement
```

inspect-expression-case:

```
inspect-pattern inspect-guardopt => assignment-expression
```

inspect-pattern:

```
wildcard-pattern
identifier-pattern
expression-pattern
structured-binding-pattern
alternative-pattern
binding-pattern
dereference-pattern
extractor-pattern
```

inspect-guard:

```
if ( expression )
```

7 Design Decisions

7.1 Extending Structured Bindings Declaration

The design is intended to be consistent and to naturally extend the notions introduced by structured bindings. That is, The subobjects are **referred** to rather than being assigned into new variables.

7.2 inspect rather than switch

This proposal introduces a new `inspect` statement rather than trying to extend the `switch` statement. P0095R0 [4] had proposed extending `switch` and received feedback to “leave `switch` alone” in Kona 2015.

The following are some of the reasons considered:

- `switch` allows the `case` labels to appear **anywhere**, which hinders the goal of pattern matching in providing **structured** inspection.
- The fall-through semantics of `switch` generally results in `break` being attached to every case, and is known to be error-prone.
- `switch` is purposely restricted to integrals for **guaranteed** efficiency. The primary goal of pattern matching in this paper is expressiveness while being at least as efficient as the naively hand-written code.

7.3 First Match rather than Best Match

The proposed matching algorithm has first match semantics. The choice of first match is mainly due to complexity. Our overload resolution rules for function declarations are extremely complex and is often a mystery.

Best match via overload resolution for function declarations are absolutely necessary due to the non-local and unordered nature of declarations. That is, function declarations live in different files and get pulled in via mechanisms such as `#include` and `using` declarations, and there is no defined order of declarations like Haskell does, for example. If function dispatching depended on the order of `#include` and/or `using` declarations being pulled in from hundreds of files, it would be a complete disaster.

Pattern matching on the other hand do not have this problem because the construct is local and ordered in nature. That is, all of the candidate patterns appear locally within `inspect (x) { /* ... */ }` which cannot span across multiple files, and appear in a specified order. This is consistent with `try/catch` for the same reasons: locality and order.

Consider also the amount of limitations we face in overload resolution due to the opacity of user-defined types. `T*` is related to `unique_ptr<T>` as it is to `vector<T>` as far as the type system is concerned. This limitation will likely be even bigger in a pattern matching context with the amount of customization points available for user-defined behavior.

7.4 Language rather than Library

There are three popular pattern matching libraries for C++ today: Mach7 [6], MPark.Patterns [3], and `simple_match` [1].

While the libraries have been useful for gaining experience with interfaces and implementation, the issue of introducing identifiers, syntactic overhead of the patterns, and the reduced optimization opportunities justify support as a language feature from a usability standpoint.

7.5 Matchers and Extractors

Many languages provide a wide array of patterns through various syntactic forms. While this is a potential direction for C++, it would mean that every new type of matching requires new syntax to be added to the language. This would result in a narrow set of types being supported through limited customization points.

Matchers and extractors are supported in order to minimize the number of patterns with special syntax. The following are example matchers and extractors that commonly have special syntax in other languages.

Matchers / Extractors	Other Languages
<code>any_of{1, 2, 3}</code>	<code>1 2 3</code>
<code>within{1, 10}</code>	<code>1..10</code>
<code>(both! [[x, 0], [0, y]])</code>	<code>[x, 0] & [0, y]</code>
<code>(at! [p, [x, y]])</code>	<code>p @ [x, y]</code>

Each of the matchers and extractors can be found in the [Examples](#) section.

7.6 Wildcard Syntax

P1110 [10] “A placeholder with no name” discusses wildcard/placeholder syntax in different contexts of the language. The intent of this proposal is to consolidate with the results of the decision of P1110 [10].

8 Runtime Performance

The following are few of the optimizations that are worth noting.

8.0.1 Structured Binding Pattern

Structured binding patterns can be optimized by performing `switch` over the columns with the duplicates removed, rather than the naive approach of performing a comparison per element. This removes unnecessary duplicate comparisons that would be performed otherwise. This would likely require some wording around “comparison elision” in order to enable such optimizations.

8.0.2 Alternative Pattern

The sequence of alternative patterns can be executed in a `switch`.

8.0.3 Open Class Hierarchy

N3449 [8] describes techniques involving vtable pointer caching and hash conflict minimization that are implemented in the Mach7 [6] library, but also mentions further opportunities available for a compiler solution.

9 Examples

9.1 Predicate-based Discriminator

Short-string optimization using a **predicate** as a discriminator rather than an explicitly stored **value**. Adapted from the pattern matching presentation given by Bjarne Stroustrup at Urbana-Champaign 2015.

```

struct String {
    enum Storage { Local, Remote };

    int size;
    union {
        char local[32];
        struct { char *ptr; int unused_allocated_space; } remote;
    };

    // Predicate-based discriminator derived from `size`.
    Storage index() const { return size > sizeof(local) ? Remote : Local; }

    // Opt into Variant-Like protocol.
    template <Storage S>
    auto &&get() {
        if constexpr (S == Local) return local;
        else if constexpr (S == Remote) return remote;
    }

    char *data();
};

namespace std {
    // Opt into Variant-Like protocol.

    template <>
    struct variant_size<String> : std::integral_constant<std::size_t, 2> {};

    template <>
    struct variant_alternative<String::Local, String> {
        using type = decltype(String::local);
    };

    template <>
    struct variant_alternative<String::Remote, String> {
        using type = decltype(String::remote);
    };
}

char* String::data() {
    inspect (*this) {
        <Local> l: return l;
        <Remote> r: return r.ptr;
    }
    // switch (index()) {
    //     case Local: {
    //         std::variant_alternative_t<Local, String> &l = get<Local>();
    //         return l;
    //     }
    //     case Remote: {
    //         std::variant_alternative_t<Remote, String> &r = get<Remote>();

```

```

//      return r.ptr;
//    }
// }
}

```

9.2 “Closed” Class Hierarchy

A class hierarchy can effectively be closed with an `enum` that maintains the list of its members, and provide efficient dispatching by opting into the Variant-Like protocol.

A generalized mechanism of pattern is used extensively in LLVM; `llvm/Support/YAMLParse.h` [11] is an example.

```

struct Shape { enum Kind { Circle, Rectangle } kind; };

struct Circle : Shape {
    Circle(int radius) : Shape{Shape::Kind::Circle}, radius(radius) {}

    int radius;
};

struct Rectangle : Shape {
    Rectangle(int width, int height)
        : Shape{Shape::Kind::Rectangle}, width(width), height(height) {}

    int width, height;
};

namespace std {
    template <>
    struct variant_size<Shape> : std::integral_constant<std::size_t, 2> {};

    template <>
    struct variant_alternative<Shape::Circle, Shape> { using type = Circle; };

    template <>
    struct variant_alternative<Shape::Rectangle, Shape> { using type = Rectangle; };
}

Shape::Kind index(const Shape& shape) { return shape.kind; }

template <Kind K>
auto&& get(const Shape& shape) {
    return static_cast<const std::variant_alternative_t<K, Shape>&>(shape);
}

int get_area(const Shape& shape) {
    inspect (shape) {
        <Circle> c: return 3.14 * c.radius * c.radius;
        <Rectangle> r: return r.width * r.height;
    }
}

```

```

// switch (index(shape)) {
//     case Shape::Circle: {
//         const std::variant_alternative_t<Shape::Circle, Shape>& c =
//             get<Shape::Circle>(shape);
//         return 3.14 * c.radius * c.radius;
//     }
//     case Shape::Rectangle: {
//         const std::variant_alternative_t<Shape::Rectangle, Shape>& r =
//             get<Shape::Rectangle>(shape);
//         return r.width * r.height;
//     }
// }
}

```

9.3 Matcher: any_of

The logical-or pattern in other languages is typically spelled *pattern₀ | pattern₁ | ... | pattern_N*, and matches value *v* if any *pattern_i* matches *v*.

This provides a restricted form (constant-only) of the logical-or pattern.

```

template <typename... Ts>
struct any_of : std::tuple<Ts...> {
    using tuple::tuple;

    template <typename U>
    bool match(const U& u) const {
        return std::apply([&](const auto&... xs) {
            return (... || xs == u);
        }, *this);
    }
};

```

```

int fib(int n) {
    inspect (n) {
        x if (x < 0): return 0;
        ~(any_of{1, 2}): return n; // 1 | 2
        x: return fib(x - 1) + fib(x - 2);
    }
}

```

9.4 Matcher: within

The range pattern in other languages is typically spelled *first..last*, and matches *v* if *v* ∈ [*first*, *last*].

```

struct within {
    int first, last;

    bool match(int n) const { return first <= n && n <= last; }
};

```

```
inspect (n) {
  ^(within{1, 10}): { // 1..10
    std::cout << n << " is in [1, 10].";
  }
  -: {
    std::cout << n << " is not in [1, 10].";
  }
}
```

9.5 Extractor: both

The logical-and pattern in other languages is typically spelled $pattern_0 \& pattern_1 \& \dots \& pattern_N$, and matches v if all of $pattern_i$ matches v .

This extractor emulates binary logical-and with a `std::pair` where both elements are references to value v .

```
inline constexpr struct {
  template <typename U>
  std::pair<U&&, U&&> extract(U&& u) const {
    return {std::forward<U>(u), std::forward<U>(u)};
  }
} both;
```

```
inspect (v) {
  (both! [[x, 0], [0, y]]): // ...
}
```

9.6 Extractor: at

The binding pattern in other languages is typically spelled $identifier @ pattern$, binds $identifier$ to v and matches if $pattern$ matches v . This is a special case of the logical-and pattern ($_pattern_0 \& pattern_1$) where $pattern_0$ is an *identifier*. That is, $identifier \& pattern$ has the same semantics as $identifier @ pattern$, which means we get `at` for free from `both` above.

```
inline constexpr at = both;
```

```
inspect (v) {
  <Point> (at! [p, [x, y]]): // ...
  // ...
}
```

10 Future Work

10.1 Language Support for Variant

The design of this proposal also accounts for a potential language support for variant. It achieves this by keeping the alternative pattern flexible for new extensions via $\< new_entity \> pattern$.

Consider an extension to `union` that allows it to be tagged by an integral, and has proper lifetime management such that the active alternative need not be destroyed manually.

```
// `: type` specifies the type of the underlying tag value.
union U : int { char small[32]; std::vector<char> big; };
```

We could then allow `< qualified-id >` that refers to a `union` alternative to support pattern matching.

```
U u = /* ... */;

inspect (u) {
    <U::small> s: std::cout << s;
    <U::big> b: std::cout << b;
}
```

The main point is that whatever entity is introduced as the discriminator, the presented form of alternative pattern should be extendable to support it.

10.2 Patterns in range-based for loop

```
for (auto&& [0, y] : points) {
    // only operate on points on the y-axis.
}
```

Structured binding declaration is allowed in range-based `for` loop:

```
for (auto&& [x, y] : points) { /* ... */ }
```

The `[x, y]` part can also be a pattern of an `inspect` statement rather than a structured binding declaration.

Before	After
<pre>for (auto&& p : points) { auto&& [x, y] = p; // ... }</pre>	<pre>for (auto&& p : points) { inspect (p) { [x, y]: // ... } }</pre>

With this model, allowing patterns directly in range-based `for` loop becomes natural.

Code	Expanded
<pre>for (auto&& [0, y] : points) { // only points on the y-axis. }</pre>	<pre>for (auto&& p : points) { inspect (p) { [0, y]: // ... } // falls through if no match }</pre>

10.3 Note on Ranges

The benefit of pattern matching for ranges is unclear. While it's possible to come up with a ranges pattern, e.g., `{x, y, z}` to match against a fixed-size range, it's not clear whether there is a worthwhile benefit.

The typical pattern found in functional languages of matching a range on head and tail doesn't seem to be all that common or useful in C++ since ranges are generally handled via loops rather than recursion.

Ranges likely will be best served by the range adaptors / algorithms, but further investigation is needed.

11 Acknowledgements

Thanks to all of the following:

- Yuriy Solodkyy, Gabriel Dos Reis, Bjarne Stroustrup for their prior work on N3449 [8], Open Pattern Matching for C++ [7], and the Mach7 [6] library.
- Pattern matching presentation by Bjarne Stroustrup at Urbana-Champaign 2015.
- David Sankel for his work on P0095 [5].
- Jeffrey Yasskin/JF Bastien for their work on P1110 [10].
- (In alphabetical order by last name) John Bandela, Agustín Bergé, Ori Bernstein, Matt Calabrese, Alexander Chow, Louis Dionne, Michał Dominiak, Eric Fiselier, Zach Laine, Jason Lucas, David Sankel, Bjarne Stroustrup, Tony Van Eerd, and everyone else who contributed to the discussions, and encouraged me to write this paper.

References

- [1] John Bandela. Simple, Extensible C++ Pattern Matching Library. Retrieved from https://github.com/jbandela/simple_match
- [2] Luc Maranget. Warnings for pattern matching. Retrieved from <http://moscova.inria.fr/~maranget/papers/warn/index.html>
- [3] Michael Park. Pattern Matching in C++. *MPark.Patterns*. Retrieved from <https://github.com/mpark/patterns>
- [4] David Sankel. 2015. Pattern Matching and Language Variants. *P0095R0*. Retrieved from <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/p0095r0.html>
- [5] David Sankel. 2016. Pattern Matching and Language Variants. *P0095R1*. Retrieved from <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2016/p0095r1.html>
- [6] Yuriy Solodkyy, Gabriel Dos Reis, and Bjarne Stroustrup. Mach7: Pattern Matching for C++. *Mach7*. Retrieved from <https://github.com/solodon4/Mach7>
- [7] Yuriy Solodkyy, Gabriel Dos Reis, and Bjarne Stroustrup. Open Pattern Matching for C++. Retrieved from <http://www.stroustrup.com/OpenPatternMatching.pdf>
- [8] Bjarne Stroustrup. 2012. Open and Efficient Type Switch for C++. *N3449*. Retrieved from <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2012/n3449.pdf>
- [9] Herb Sutter, Bjarne Stroustrup, and Gabriel Dos Reis. 2016. Structured bindings. *P0144*. Retrieved from <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2016/p0144r2.pdf>

- [10] Jeffrey Yasskin and JF Bastien. A placeholder with no name. Retrieved from <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2018/p1110r0.html>
- [11] Retrieved from http://llvm.org/doxygen/YAMLParse_8h_source.html